

Федеральное государственное автономное  
образовательное учреждение  
высшего образования  
«СИБИРСКИЙ ФЕДЕРАЛЬНЫЙ УНИВЕРСИТЕТ»

Институт космических и информационных технологий  
институт

Кафедра «Информатика»  
кафедра

**ОТЧЕТ О ПРАКТИЧЕСКОЙ РАБОТЕ № 8**

Методы оптимизации  
Тема

Преподаватель		Е. П. Моргунов
	Подпись, дата	Инициалы, Фамилия
Студент	КИ19-17/1Б, №031939174	А. К. Никитин
	Номер группы, зачетной книжки	Инициалы, Фамилия
	Подпись, дата	

Красноярск 2021

## 1 Ход выполнения

### 1.1 Вопрос 1

Для того, чтобы найти в индексе первую строку в соответствии с требуемым порядком, нужно некоторое время.

### 1.2 Вопрос 3

```
demo=# EXPLAIN WITH some_seats AS (  
demo(# SELECT * FROM seats  
demo(# WHERE aircraft_code = '319'  
demo(# )  
demo-# SELECT * FROM some_seats  
demo-# WHERE fare_conditions = 'Business';  
                                QUERY PLAN  
-----  
CTE Scan on some_seats (cost=10.43..13.04 rows=1 width=74)  
  Filter: ((fare_conditions)::text = 'Business'::text)  
    CTE some_seats  
      -> Index Scan using seats_aircraft_code_idx on seats (cost=0.28..10.43 rows=116 width=15)  
          Index Cond: (aircraft_code = '319'::bpchar)  
(5 строк)
```

Рисунок 1 – Работа планировщика в случае общих табличных выражений

Как мы видим, планировщик тратит время только на поиск внутри табличного выражения, при этом по самому табличному выражению он использует CTE Scan.

### 1.3 Вопрос 5

Время 3.04 – это время, которое потребуется непосредственно на фильтрацию сгруппированной таблицы, на остальные операции потребуется время, поэтому первая оценка стоимости равна 3.82.

## 1.4 Вопрос 7

```
demo=# BEGIN;
BEGIN
demo=# EXPLAIN
demo-# INSERT INTO seats
demo-# VALUES
demo-# (320, '10G', 'Economy'),
demo-# (320, '10H', 'Business');
                                QUERY PLAN
-----
Insert on seats  (cost=0.00..0.03 rows=2 width=74)
->  Values Scan on "*VALUES*" (cost=0.00..0.03 rows=2 width=74)
(2 строки)
```

Рисунок 2 – Вставка строк

```
demo=# BEGIN;
BEGIN
demo=# EXPLAIN
demo-# UPDATE aircrafts
demo-# SET range = range + 100
demo-# WHERE model ~ '^Аэро';
                                QUERY PLAN
-----
Update on aircrafts_data ml  (cost=0.00..3.39 rows=1 width=58)
->  Seq Scan on aircrafts_data ml  (cost=0.00..3.39 rows=1 width=58)
      Filter: ((model ->> lang()) ~ '^Аэро'::text)
(3 строки)
```

Рисунок 3 – Обновление строк таблицы

```
demo=# BEGIN;
BEGIN
demo=# EXPLAIN
demo-# DELETE FROM aircrafts
demo-# WHERE model ~ '^Боинг';
                                QUERY PLAN
-----
Delete on aircrafts_data ml  (cost=0.00..3.39 rows=1 width=6)
->  Seq Scan on aircrafts_data ml  (cost=0.00..3.39 rows=1 width=6)
      Filter: ((model ->> lang()) ~ '^Боинг'::text)
(3 строки)
```

Рисунок 4 – Удаление строк таблицы

## 1.5 Вопрос 9

```
demo=# EXPLAIN ANALYZE
demo=# SELECT * FROM routes;

QUERY PLAN

-----
Hash Join (cost=2705.54..3008.86 rows=276 width=252) (actual time=126.502..142.045 rows=710 loops=1)
  Hash Cond: (f3.arrival_airport = ml_1.airport_code)
  CTE f3
    -> GroupAggregate (cost=2401.55..2694.86 rows=1020 width=67) (actual time=108.451..112.718 rows=710 loops=1)
      Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure))
      -> Sort (cost=2401.55..2427.06 rows=10202 width=39) (actual time=108.438..109.332 rows=3798 loops=1)
        Sort Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure)), ((to_char(flights.scheduled_departure, 'ID'::text))::integer)
        Sort Method: quicksort Memory: 332kB
      -> HashAggregate (cost=1518.24..1722.28 rows=10202 width=39) (actual time=94.101..96.264 rows=3798 loops=1)
        Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, (flights.scheduled_arrival - flights.scheduled_departure), (to_char(flights.scheduled_departure, 'ID'::text))::integer
        -> Seq Scan on flights (cost=0.00..1021.42 rows=33121 width=39) (actual time=10.504..58.487 rows=33121 loops=1)
  -> Hash Join (cost=5.34..28.47 rows=530 width=234) (actual time=108.577..114.316 rows=710 loops=1)
    Hash Cond: (f3.departure_airport = ml_1.airport_code)
    -> CTE Scan on f3 (cost=0.00..20.40 rows=1020 width=124) (actual time=108.454..113.547 rows=710 loops=1)
    -> Hash (cost=4.04..4.04 rows=104 width=114) (actual time=0.113..0.113 rows=104 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 18kB
      -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=114) (actual time=0.012..0.053 rows=104 loops=1)
  -> Hash (cost=4.04..4.04 rows=104 width=114) (actual time=17.674..17.674 rows=104 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 18kB
    -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=114) (actual time=16.973..17.603 rows=104 loops=1)
Planning Time: 121.172 ms
Execution Time: 143.285 ms
(22 строки)
```

Рисунок 5 – Анализ выборки из готового представления

```
QUERY PLAN

-----
Hash Join (cost=2705.54..3008.86 rows=276 width=252) (actual time=87.809..104.541 rows=710 loops=1)
  Hash Cond: (f3.arrival_airport = ml_1.airport_code)
  CTE f3
    -> GroupAggregate (cost=2401.55..2694.86 rows=1020 width=67) (actual time=87.545..92.061 rows=710 loops=1)
      Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure))
      -> Sort (cost=2401.55..2427.06 rows=10202 width=39) (actual time=87.531..88.454 rows=3798 loops=1)
        Sort Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure)), ((to_char(flights.scheduled_departure, 'ID'::text))::integer)
        Sort Method: quicksort Memory: 332kB
      -> HashAggregate (cost=1518.24..1722.28 rows=10202 width=39) (actual time=74.265..76.061 rows=3798 loops=1)
        Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, (flights.scheduled_arrival - flights.scheduled_departure), (to_char(flights.scheduled_departure, 'ID'::text))::integer
        -> Seq Scan on flights (cost=0.00..1021.42 rows=33121 width=39) (actual time=0.012..40.119 rows=33121 loops=1)
  -> Hash Join (cost=5.34..28.47 rows=530 width=234) (actual time=87.665..93.596 rows=710 loops=1)
    Hash Cond: (f3.departure_airport = ml_1.airport_code)
    -> CTE Scan on f3 (cost=0.00..20.40 rows=1020 width=124) (actual time=87.549..92.810 rows=710 loops=1)
    -> Hash (cost=4.04..4.04 rows=104 width=114) (actual time=0.107..0.107 rows=104 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 18kB
      -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=114) (actual time=0.010..0.051 rows=104 loops=1)
  -> Hash (cost=4.04..4.04 rows=104 width=114) (actual time=0.096..0.096 rows=104 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 18kB
    -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=114) (actual time=0.008..0.047 rows=104 loops=1)
Planning Time: 0.346 ms
Execution Time: 143.236 ms
(22 строки)
```

Рисунок 6 – Анализ формирующей выборки

## 1.6 Вопрос 11

```
demo=# EXPLAIN ANALYZE
demo=# SELECT * FROM flights_v;

QUERY PLAN

-----
Hash Join (cost=10.68..34830.88 rows=33121 width=255) (actual time=0.242..625.935 rows=33121 loops=1)
  Hash Cond: (f.arrival_airport = ml_1.airport_code)
  -> Hash Join (cost=5.34..786.03 rows=33121 width=188) (actual time=0.099..47.708 rows=33121 loops=1)
    Hash Cond: (f.departure_airport = ml_1.airport_code)
    -> Seq Scan on flights f (cost=0.00..690.21 rows=33121 width=63) (actual time=0.004..11.029 rows=33121 loops=1)
    -> Hash (cost=4.04..4.04 rows=104 width=129) (actual time=0.090..0.090 rows=104 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 20kB
      -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=129) (actual time=0.002..0.041 rows=104 loops=1)
  -> Hash (cost=4.04..4.04 rows=104 width=129) (actual time=0.099..0.099 rows=104 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 20kB
    -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=129) (actual time=0.007..0.047 rows=104 loops=1)
Planning Time: 0.427 ms
Execution Time: 633.153 ms
(13 строк)

demo=# EXPLAIN ANALYZE
demo=# SELECT * FROM flights_tt;

QUERY PLAN

-----
Seq Scan on flights tt (cost=0.00..1027.20 rows=17120 width=362) (actual time=0.006..10.341 rows=33121 loops=1)
Planning Time: 0.123 ms
Execution Time: 16.340 ms
(3 строки)
```

## Рисунок 7 – Выполнение простых запросов для исходной и временной таблицы

Так как таблица `flights_v` является представлением, то она формируется каждый раз при обращении к ней. Временная же таблица формируется лишь один раз. Отсюда и разница во времени выполнения запроса и Query plan.

```
demo=# EXPLAIN ANALYZE SELECT arrival_city, count ( * ) FROM flights_v
demo=# WHERE arrival_city = 'Yfa'
demo=# GROUP BY arrival_city;

QUERY PLAN
-----
GroupAggregate (cost=35.91..899.62 rows=1 width=40) (actual time=22.662..22.662 rows=1 loops=1)
  Group Key: (ml_1.city ->> lang())
  -> Hash Join (cost=35.91..897.77 rows=318 width=32) (actual time=0.475..22.508 rows=384 loops=1)
    Hash Cond: (f.departure_airport = ml.airport_code)
    -> Hash Join (cost=30.57..811.26 rows=318 width=53) (actual time=0.387..20.900 rows=384 loops=1)
      Hash Cond: (f.arrival_airport = ml_1.airport_code)
      -> Seq Scan on flights f (cost=0.00..690.21 rows=33121 width=8) (actual time=0.004..9.093 rows=33121 loops=1)
      -> Hash (cost=30.56..30.56 rows=1 width=53) (actual time=0.376..0.376 rows=1 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 5kB
        -> Seq Scan on airports_data ml_1 (cost=0.00..30.56 rows=1 width=53) (actual time=0.301..0.374 rows=1 loops=1)
          Filter: ((city ->> lang()) = 'Yfa'::text)
          Rows Removed by Filter: 103
    -> Hash (cost=4.04..4.04 rows=104 width=4) (actual time=0.080..0.080 rows=104 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 7kB
      -> Seq Scan on airports_data ml (cost=0.00..4.04 rows=104 width=4) (actual time=0.006..0.042 rows=104 loops=1)
Planning Time: 0.379 ms
Execution Time: 22.709 ms
(17 строк)

demo=# EXPLAIN ANALYZE SELECT arrival_city, count ( * ) FROM flights_tt
demo=# WHERE arrival_city = 'Yfa'
demo=# GROUP BY arrival_city;

QUERY PLAN
-----
GroupAggregate (cost=0.00..1071.13 rows=70 width=40) (actual time=10.550..10.551 rows=1 loops=1)
  Group Key: arrival_city
  -> Seq Scan on flights_tt (cost=0.00..1070.00 rows=86 width=32) (actual time=0.012..10.428 rows=384 loops=1)
    Filter: (arrival_city = 'Yfa'::text)
    Rows Removed by Filter: 32737
Planning Time: 0.146 ms
Execution Time: 10.577 ms
(7 строк)
```

## Рисунок 8 – Сложный запрос для представления и временной таблицы

### 1.7 Вопрос 13

```
GroupAggregate (cost=13325.89..13337.75 rows=200 width=16) (actual time=165.432..165.432 rows=0 loops=1)
  Group Key: count_tickets.num_tickets
  -> Sort (cost=13325.89..13329.18 rows=1314 width=8) (actual time=165.431..165.431 rows=0 loops=1)
    Sort Key: count_tickets.num_tickets DESC
    Sort Method: quicksort Memory: 17kB
    -> Subquery Scan on count_tickets (cost=13034.17..13257.83 rows=1314 width=8) (actual time=165.428..165.429 rows=0 loops=1)
      -> Finalize GroupAggregate (cost=13034.17..13244.69 rows=1314 width=15) (actual time=165.427..165.427 rows=0 loops=1)
        Group Key: b.book_ref
        -> Gather Merge (cost=13034.17..13223.91 rows=1528 width=15) (actual time=165.427..165.921 rows=0 loops=1)
          Workers Planned: 2
          Workers Launched: 2
          -> Partial GroupAggregate (cost=12034.14..12047.51 rows=764 width=15) (actual time=156.379..156.379 rows=0 loops=3)
            Group Key: b.book_ref
            -> Sort (cost=12034.14..12036.05 rows=764 width=7) (actual time=156.378..156.378 rows=0 loops=3)
              Sort Key: b.book_ref
              Sort Method: quicksort Memory: 17kB
              Worker 0: Sort Method: quicksort Memory: 17kB
              Worker 1: Sort Method: quicksort Memory: 17kB
              -> Parallel Hash Join (cost=4002.38..11997.56 rows=764 width=7) (actual time=156.314..156.314 rows=0 loops=3)
                Hash Cond: (t.book_ref = b.book_ref)
                -> Parallel Seq Scan on tickets t (cost=0.00..7594.05 rows=152805 width=7) (never executed)
                -> Parallel Hash (cost=3992.72..3992.72 rows=773 width=7) (actual time=156.255..156.256 rows=0 loops=3)
                  Buckets: 2048 Batches: 1 Memory Usage: 0kB
                  -> Parallel Seq Scan on bookings b (cost=0.00..3992.72 rows=773 width=7) (actual time=154.382..154.382 rows=0 loops=3)
                    Filter: (date_trunc('mon'::text, book_date) = '2016-09-01 00:00:00+07'::timestamp with time zone)
                    Rows Removed by Filter: 87596
Planning Time: 0.233 ms
Execution Time: 165.998 ms
```

## Рисунок 9 – Исходный запрос

```

QUERY PLAN
-----
GroupAggregate (cost=30592.96..30604.82 rows=200 width=16) (actual time=1975.845..1975.845 rows=0 loops=1)
  Group Key: count_tickets.num_tickets
  -> Sort (cost=30592.96..30596.25 rows=1314 width=8) (actual time=1975.844..1975.844 rows=0 loops=1)
    Sort Key: count_tickets.num_tickets DESC
    Sort Method: quicksort Memory: 17kB
    -> Subquery Scan on count_tickets (cost=29525.06..30524.90 rows=1314 width=8) (actual time=1975.841..1975.841 rows=0 loops=1)
      -> Finalize GroupAggregate (cost=29525.06..30511.76 rows=1314 width=15) (actual time=1975.840..1975.840 rows=0 loops=1)
        Group Key: b.book_ref
        -> Gather Merge (cost=29525.06..30490.98 rows=1528 width=15) (actual time=1975.839..1975.866 rows=0 loops=1)
          Workers Planned: 2
          Workers Launched: 2
          -> Partial GroupAggregate (cost=28525.04..29314.59 rows=764 width=15) (actual time=1951.051..1951.051 rows=0 loops=3)
            Group Key: b.book_ref
            -> Merge Join (cost=28525.04..29303.13 rows=764 width=7) (actual time=1951.050..1951.050 rows=0 loops=3)
              Merge Cond: (t.book_ref = b.book_ref)
              -> Sort (cost=22841.08..23223.10 rows=152805 width=7) (actual time=1505.689..1505.689 rows=1 loops=3)
                Sort Key: t.book_ref
                Sort Method: external merge Disk: 2136kB
                Worker 0: Sort Method: external merge Disk: 1880kB
                Worker 1: Sort Method: external merge Disk: 2104kB
                -> Parallel Seq Scan on tickets t (cost=0.00..7594.05 rows=152805 width=7) (actual time=266.865..758.698 rows=122244 loops=3)
              -> Sort (cost=5683.88..5687.17 rows=1314 width=7) (actual time=445.358..445.358 rows=0 loops=3)
                Sort Key: b.book_ref
                Sort Method: quicksort Memory: 17kB
                Worker 0: Sort Method: quicksort Memory: 17kB
                Worker 1: Sort Method: quicksort Memory: 17kB
                -> Seq Scan on bookings b (cost=0.00..5615.82 rows=1314 width=7) (actual time=445.351..445.351 rows=0 loops=3)
                  Filter: (date_trunc('mon':text, book_date) = '2016-09-01 00:00:00+07':timestamp with time zone)
                  Rows Removed by Filter: 262788
Planning Time: 0.252 ms
Execution Time: 1976.227 ms
(31 строка)

```

Рисунок 10 – Запрос с отключенным hashjoin

```

QUERY PLAN
-----
GroupAggregate (cost=30592.96..30604.82 rows=200 width=16) (actual time=1377.964..1377.965 rows=0 loops=1)
  Group Key: count_tickets.num_tickets
  -> Sort (cost=30592.96..30596.25 rows=1314 width=8) (actual time=1377.964..1377.964 rows=0 loops=1)
    Sort Key: count_tickets.num_tickets DESC
    Sort Method: quicksort Memory: 17kB
    -> Subquery Scan on count_tickets (cost=29525.06..30524.90 rows=1314 width=8) (actual time=1377.961..1377.961 rows=0 loops=1)
      -> Finalize GroupAggregate (cost=29525.06..30511.76 rows=1314 width=15) (actual time=1377.960..1377.960 rows=0 loops=1)
        Group Key: b.book_ref
        -> Gather Merge (cost=29525.06..30490.98 rows=1528 width=15) (actual time=1377.959..1377.988 rows=0 loops=1)
          Workers Planned: 2
          Workers Launched: 2
          -> Partial GroupAggregate (cost=28525.04..29314.59 rows=764 width=15) (actual time=1359.166..1359.166 rows=0 loops=3)
            Group Key: b.book_ref
            -> Merge Join (cost=28525.04..29303.13 rows=764 width=7) (actual time=1359.164..1359.164 rows=0 loops=3)
              Merge Cond: (t.book_ref = b.book_ref)
              -> Sort (cost=22841.08..23223.10 rows=152805 width=7) (actual time=899.912..899.912 rows=1 loops=3)
                Sort Key: t.book_ref
                Sort Method: external merge Disk: 2120kB
                Worker 0: Sort Method: external merge Disk: 1968kB
                Worker 1: Sort Method: external merge Disk: 2024kB
                -> Parallel Seq Scan on tickets t (cost=0.00..7594.05 rows=152805 width=7) (actual time=0.020..163.917 rows=122244 loops=3)
              -> Sort (cost=5683.88..5687.17 rows=1314 width=7) (actual time=459.249..459.249 rows=0 loops=3)
                Sort Key: b.book_ref
                Sort Method: quicksort Memory: 17kB
                Worker 0: Sort Method: quicksort Memory: 17kB
                Worker 1: Sort Method: quicksort Memory: 17kB
                -> Seq Scan on bookings b (cost=0.00..5615.82 rows=1314 width=7) (actual time=459.242..459.245 rows=0 loops=3)
                  Filter: (date_trunc('mon':text, book_date) = '2016-09-01 00:00:00+07':timestamp with time zone)
                  Rows Removed by Filter: 262788
Planning Time: 0.223 ms
Execution Time: 1378.353 ms
(31 строка)

```

Рисунок 11 – Запрос с отключенным nested loop

Везде используется Seq scan.

## 1.8 Вопрос 15

```

demo=# EXPLAIN SELECT model, range,
demo=# CASE WHEN range < 2000 THEN 'Ближнемагистральный'
demo=# WHEN range < 5000 THEN 'Среднемагистральный'
demo=# ELSE 'Дальнемагистральный'
demo=# END AS type
demo=# FROM aircrafts
demo=# ORDER BY model;

```

```

QUERY PLAN
-----
Sort (cost=3.55..3.57 rows=9 width=68)
  Sort Key: ((ml.model ->> lang()))
  -> Seq Scan on aircrafts_data ml (cost=0.00..3.41 rows=9 width=68)
(3 строки)

```



## Рисунок 12 – Первый запрос

```
demo=# EXPLAIN SELECT s.seat_no, s.fare_conditions
FROM seats s
JOIN aircrafts a ON s.aircraft_code = a.aircraft_code
WHERE a.model ~ '^Cessna'
ORDER BY s.seat_no;

QUERY PLAN

-----
Sort (cost=21.26..21.63 rows=149 width=11)
  Sort Key: s.seat_no
  -> Nested Loop (cost=0.28..15.88 rows=149 width=11)
    -> Seq Scan on aircrafts_data ml (cost=0.00..3.39 rows=1 width=16)
        Filter: ((model ->> lang()) ~ '^Cessna'::text)
    -> Index Scan using seats_aircraft_code_idx on seats s (cost=0.28..11.00 rows=149 width=15)
        Index Cond: (aircraft_code = ml.aircraft_code)
(7 строк)
```

## Рисунок 13 – Второй запрос

```
demo=# EXPLAIN SELECT a.aircraft_code AS a_code,
demo-# a.model,
demo-# r.aircraft_code AS r_code,
demo-# count( r.aircraft_code ) AS num_routes
demo-# FROM aircrafts a
demo-# LEFT OUTER JOIN routes r ON r.aircraft_code = a.aircraft_code
demo-# GROUP BY 1, 2, 3
demo-# ORDER BY 4 DESC;

QUERY PLAN

-----
Sort (cost=2741.55..2741.58 rows=12 width=72)
  Sort Key: (count(f3.aircraft_code)) DESC
  -> GroupAggregate (cost=2738.04..2741.34 rows=12 width=72)
    Group Key: ml.aircraft_code, ((ml.model ->> lang())), f3.aircraft_code
    -> Sort (cost=2738.04..2738.07 rows=12 width=64)
      Sort Key: ml.aircraft_code, ((ml.model ->> lang())), f3.aircraft_code
      -> Hash Right Join (cost=2786.74..2737.82 rows=12 width=64)
        Hash Cond: (f3.aircraft_code = ml.aircraft_code)
        -> Hash Join (cost=2785.54..2738.10 rows=276 width=252)
          Hash Cond: (f3.arrival_airport = ml_2.airport_code)
          CTE f3
          -> GroupAggregate (cost=2401.55..2694.86 rows=1020 width=67)
            Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - fl
            ights.scheduled_departure))
            -> Sort (cost=2401.55..2427.06 rows=10202 width=39)
              Sort Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival
              - flights.scheduled_departure)), ((to_char(flights.scheduled_departure, 'ID'::text))::integer)
              -> HashAggregate (cost=1518.24..1722.28 rows=10202 width=39)
                Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, (flights.scheduled_ar
                rival - flights.scheduled_departure), (to_char(flights.scheduled_departure, 'ID'::text))::integer
                -> Seq Scan on flights (cost=0.00..1021.42 rows=33121 width=39)
                -> Hash Join (cost=5.34..28.47 rows=530 width=32)
                  Hash Cond: (f3.departure_airport = ml_1.airport_code)
                  -> CTE Scan on f3 (cost=0.00..20.40 rows=1020 width=48)
                  -> Hash (cost=4.04..4.04 rows=104 width=4)
                    -> Seq Scan on airports_data ml_1 (cost=0.00..4.04 rows=104 width=4)
                    -> Hash (cost=4.04..4.04 rows=104 width=4)
                      -> Seq Scan on airports_data ml_2 (cost=0.00..4.04 rows=104 width=4)
                      -> Hash (cost=1.09..1.09 rows=9 width=48)
                        -> Seq Scan on aircrafts_data ml (cost=0.00..1.09 rows=9 width=48)
(27 строк)
```

## Рисунок 14 – Третий запрос

```
demo=# EXPLAIN SELECT count( * )
demo-# FROM ( ticket_flights t
demo-# JOIN flights f ON t.flight_id = f.flight_id
demo-# )
demo-# LEFT OUTER JOIN boarding_passes b
demo-# ON t.ticket_no = b.ticket_no AND t.flight_id = b.flight_id
demo-# WHERE f.actual_departure IS NOT NULL AND b.flight_id IS NULL;

QUERY PLAN

-----
Finalize Aggregate (cost=33693.92..33693.93 rows=1 width=8)
  -> Gather (cost=33693.70..33693.91 rows=2 width=8)
    Workers Planned: 2
    -> Partial Aggregate (cost=32693.70..32693.71 rows=1 width=8)
      -> Parallel Hash Anti Join (cost=12234.47..32353.55 rows=136063 width=0)
        Hash Cond: ((t.ticket_no = b.ticket_no) AND (t.flight_id = b.flight_id))
        -> Hash Join (cost=900.07..14636.17 rows=220865 width=18)
          Hash Cond: (t.flight_id = f.flight_id)
          -> Parallel Seq Scan on ticket_flights t (cost=0.00..12592.19 rows=435719 width=18)
          -> Hash (cost=690.21..690.21 rows=16789 width=4)
            -> Seq Scan on flights f (cost=0.00..690.21 rows=16789 width=4)
              Filter: (actual_departure IS NOT NULL)
        -> Parallel Hash (cost=6413.36..6413.36 rows=241536 width=18)
          -> Parallel Seq Scan on boarding_passes b (cost=0.00..6413.36 rows=241536 width=18)
(14 строк)
```

## Рисунок 15 – Четвертый запрос

```

demo=# EXPLAIN SELECT arrival_city FROM routes
demo=# WHERE departure_city = 'Москва'
demo=# UNION
demo=# SELECT arrival_city FROM routes
demo=# WHERE departure_city = 'Санкт-Петербург'
demo=# ORDER BY arrival_city;
QUERY PLAN
-----
Sort (cost=5500.74..5500.76 rows=6 width=32)
  Sort Key: routes.arrival_city
  -> HashAggregate (cost=5500.60..5500.66 rows=6 width=32)
    Group Key: routes.arrival_city
    -> Append (cost=2725.57..5500.59 rows=6 width=32)
      -> Subquery Scan on routes (cost=2725.57..2750.25 rows=3 width=32)
        -> Nested Loop (cost=2725.57..2750.22 rows=3 width=252)
          CTE f3
          -> GroupAggregate (cost=2401.55..2694.86 rows=1020 width=67)
            Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure))
            -> Sort (cost=2401.55..2427.06 rows=1020 width=39)
              Sort Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, ((flights.scheduled_arrival - flights.scheduled_departure)), ((to_char(flights.scheduled_departure, 'ID'::text))::integer)
              -> HashAggregate (cost=1518.24..1722.28 rows=1020 width=39)
                Group Key: flights.flight_no, flights.departure_airport, flights.arrival_airport, flights.aircraft_code, (flights.scheduled_arrival - flights.scheduled_departure), (to_char(flights.scheduled_departure, 'ID'::text))::integer
                -> Seq Scan on flights (cost=0.00..1021.42 rows=33121 width=39)
          -> Hash Join (cost=20.57..52.71 rows=6 width=16)

```

Рисунок 16 – Пятый запрос

```

demo=# EXPLAIN SELECT min( total_amount ) FROM bookings;
QUERY PLAN
-----
Finalize Aggregate (cost=4606.38..4606.39 rows=1 width=32)
  -> Gather (cost=4606.26..4606.37 rows=1 width=32)
    Workers Planned: 1
    -> Partial Aggregate (cost=3606.26..3606.27 rows=1 width=32)
      -> Parallel Seq Scan on bookings (cost=0.00..3219.81 rows=154581 width=6)
(5 строк)

```

Рисунок 17 – Шестой запрос

```

demo=# EXPLAIN SELECT count( * ) FROM bookings
demo=# WHERE total_amount >
demo=# ( SELECT avg( total_amount ) FROM bookings );
QUERY PLAN
-----
Finalize Aggregate (cost=9341.59..9341.60 rows=1 width=8)
  InitPlan 1 (returns $1)
    -> Finalize Aggregate (cost=4606.38..4606.39 rows=1 width=32)
      -> Gather (cost=4606.27..4606.38 rows=1 width=32)
        Workers Planned: 1
        -> Partial Aggregate (cost=3606.27..3606.28 rows=1 width=32)
          -> Parallel Seq Scan on bookings bookings_1 (cost=0.00..3219.81 rows=154581 width=6)
    -> Gather (cost=4735.08..4735.19 rows=1 width=8)
      Workers Planned: 1
      Params Evaluated: $1
      -> Partial Aggregate (cost=3735.08..3735.09 rows=1 width=8)
        -> Parallel Seq Scan on bookings (cost=0.00..3606.26 rows=51527 width=0)
          Filter: (total_amount > $1)
(13 строк)

```

Рисунок 18 – Седьмой запрос

## 1.9 Вопрос 17

С разделом документации 14.2 «Статистика, используемая планировщиком» ознакомился.

## 1.10 Вопрос 19

С разделом документации 14.4 «Наполнение базы данных» ознакомился.



```
demo=# BEGIN;  
BEGIN  
demo=# COMMIT;  
COMMIT  
demo=# █
```

Рисунок 19 – Создание транзакции перед добавлением

```
demo=# BEGIN;  
BEGIN  
demo=# COPY aircrafts_temp FROM '/home/postgres/some_file';  
COPY 9  
demo=# ROLLBACK;  
ROLLBACK  
demo=# █
```

Рисунок 20 – Использование COPY вместо INSERT

```
demo=# DROP INDEX aircrafts_temp_model_idx;  
DROP INDEX  
demo=# COPY aircrafts_temp FROM '/home/postgres/some_file';  
COPY 9  
demo=# CREATE INDEX ON aircrafts_temp (model);  
CREATE INDEX
```

Рисунок 21 – Удаление индекса и его воссоздание после добавления